



Cyclistic Bike-Sharing Data Analysis

Thanakorn Thanakraikiti (Nun)

June 21, 2023

Contents

1	Introduction	3
1.1	Purpose	3
1.2	Stakeholders	3
2	Data Sources	4
2.1	Data Description	4
2.2	Data Credibility	5
3	Data Cleaning and Manipulation	6
3.1	Setting Up and Loading Data into R Environment	6
3.2	Removing Duplicate or Irrelevant Observations	7
3.3	Fixing Structural Errors	7
3.4	Handle Missing Data	9
3.5	Adding New Columns	11
3.6	Filter Unwanted Outliers	12
4	Analysis and Key Findings	13
4.1	The Number of Rides	13
4.2	The Duration of Rides	13
4.3	The Location of Rides	13
4.4	Key Findings	13
5	Conclusion and Recommendations	14
5.1	Conclusion	14
5.2	Recommendations	14

1 Introduction

Cyclistic is a successful American bicycle-sharing program that was established in 2016. Since then, the program has grown to a fleet of 5,824 bicycles that are geotracked and locked into a network of 692 stations across Chicago. The bikes can be unlocked from one station and returned to any other station in the system at any time. Cyclistic users are more likely to ride for leisure, but approximately 30% use them to commute to work each day.

Cyclistic's marketing strategy has primarily focused on building general awareness and appealing to broad consumer segments. The program offers a variety of pricing plans, including single-ride passes, full-day passes, and annual memberships. Cyclistic classifies its riders into the following groups based on these plans: *casual riders* (users who purchase single-ride passes or full-day passes) and *annual members* (users who purchase an annual membership).

Cyclistic's flexible pricing plans attract a larger customer base, but financial analysts have determined that annual members are more profitable. However, casual riders are already aware of the Cyclistic program and have chosen Cyclistic to meet their mobility needs. This suggests that a marketing campaign that targets existing customers is likely to be more effective at expanding the business than a campaign that targets only new customers.

Therefore, Cyclistic's marketing analytics team is interested in understanding how casual riders and annual members use Cyclistic bikes differently. By understanding these differences, the marketing analytics team can develop more targeted marketing strategies to convert casual riders into annual members.

1.1 Purpose

This report analyzes Cyclistic's historical bike trip data from 2022 in order to identify how annual members and casual riders use Cyclistic bikes differently. The report aims to extract insights and develop the most appropriate marketing strategies that appeal to casual riders and encourage them to subscribe to annual memberships. This will enable the business to improve productivity and respond to economic changes in a productive and efficient manner. The report outlines the behavioral differences in Cyclistic bike usage between annual members and casual riders. The report further demonstrates that the recommended solutions will help Cyclistic to expand its business growth.

1.2 Stakeholders

Key stakeholders include:

- Marketing analytics team who collect, analyze, and report data to help guide Cyclistic's marketing strategy
- Marketing manager who supports the marketing campaign to promote Cyclistic's bike-share program
- Executive team who approve the recommended marketing program

2 Data Sources

This report uses Cyclistic's historical trip data from January 2022 to December 2022 as its primary source. The data sets can be downloaded from the company's website, which can be found in the following details:

- Data Source: [Trip History Data](#)
- Data Website: [System Data Website](#)
- License: [Data License Agreement](#)

The data sets are collected by the City of Chicago (a government agency) and published by Cyclistic. They contain historical trip data for all of the company's customer (rider) data since 2013, including the following information: ride history, the location of stations for each ride, and other relevant information

2.1 Data Description

The data set is stored in twelve CSV files, one for each month of 2022 (*202201-divvy-tripdata.csv* through *202212-divvy-tripdata.csv*). Each file contains a set of fields (columns) and records (rows). Each record represents a single bike trip, which is uniquely identified by an identifier for the ride.

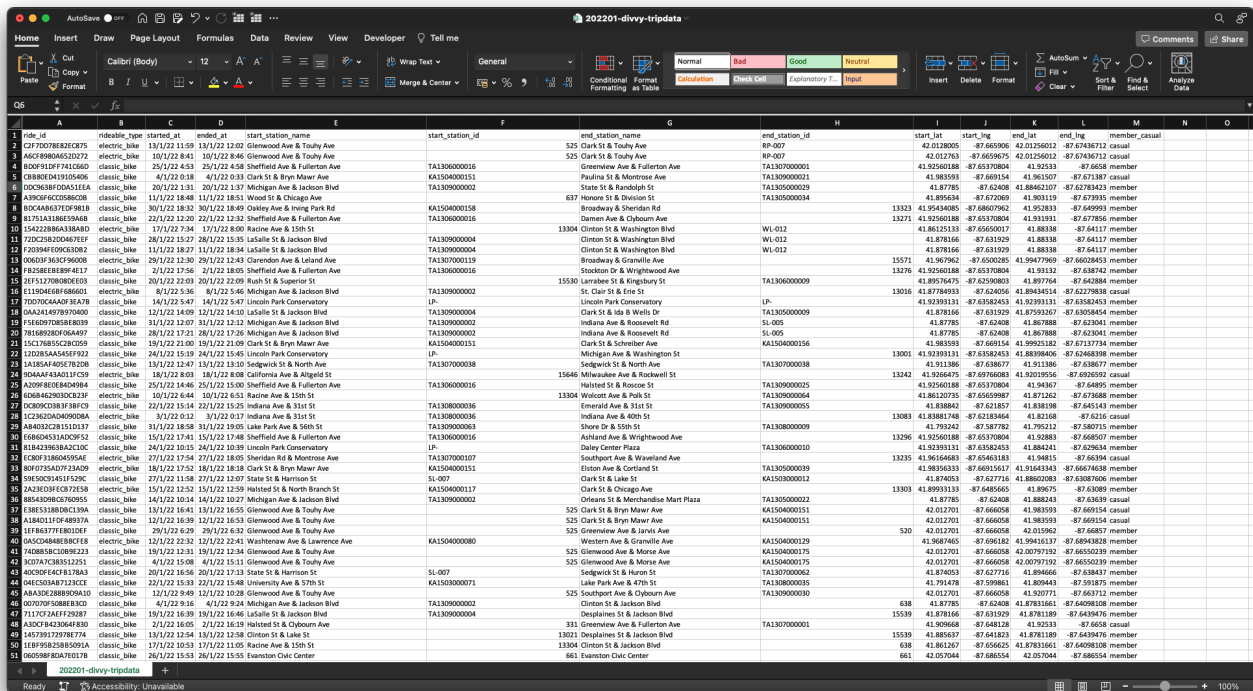


Figure 1: The spreadsheet example, containing records of Cyclistic's historical trip data in January 2022

The data sets contain the following fields:

- **ride_id** : A unique identifier for the ride
- **rideable_type** : The type of bike for the ride (classic bike, electric bike or docked bike)
- **started_at** : The date and time at which the ride started
- **ended_at** : The date and time at which the ride ended

- `start_station_id` : The ID of the station where the ride started
- `start_station_name` : The name of the station where the ride started
- `end_station_id` : The ID of the station where the ride ended
- `end_station_name` : The name of the station where the ride ended
- `start_lat` : The latitude of the station where the ride started
- `start_lng` : The longitude of the station where the ride started
- `end_lat` : The latitude of the station where the ride ended
- `end_lng` : The longitude of the station where the ride ended
- `member_casual` : The type of user who took the ride (member or casual)

2.2 Data Credibility

The data is evaluated for bias and credibility using the ROCCC method:

- **Reliable:** The data is considered reliable because it is open data provided by Cyclistic and the City of Chicago. It is a comprehensive and accurate dataset of Cyclistic's historical bike trips. The data includes information on the start and end location of each trip, the time of day the trip was taken, the type of user (casual or annual member), and the length of the trip. There are some missing values in some fields, but these can be used to explore how different customer types use Cyclistic bikes.
- **Original:** The data is considered to be original, based on a primary source of information from the company.
- **Comprehensive:** It is important to note that Cyclistic is subject to data privacy restrictions that prohibit the use of riders' personally identifiable information (PII). This means that the company cannot connect pass purchases to credit card numbers to determine if casual riders live in the Cyclistic service area or if they have purchased multiple single passes, as well as identify the rider's gender and age. Despite these restrictions, the data is complete and contains all the information needed to understand how members and casual riders use Cyclistic bikes. This includes information such as the number of rides taken, the length of each ride, the start and end stations of each ride, the time of day each ride was taken.
- **Current:** The data is current and updated monthly. The selected datasets are the most recent data available.
- **Cited:** The data is properly cited as it contains the dataset name, author, and published year.

Overall, the Cyclistic bike data is considered a credible source for analysis based on the Reliability, Originality, Comprehensiveness, Currency, and Citation (ROCCC) method. The data can be used to explore how different customer types use Cyclistic bikes. For example, the data can be used to determine which customer segments are most likely to use Cyclistic bikes for commuting, recreation, or transportation. The data can also be used to identify trends in bike usage, such as the time of day when bikes are most likely to be used or the types of bikes that are most popular.

While there are some restrictions on rider's personal information, there is still a great deal of information that can be learned from it. However, the data needs to be processed before analysis can be performed. This includes cleaning the data to remove errors and inconsistencies, and transforming the data into a format that is suitable for analysis. The data processing will be performed in the next section.

3 Data Cleaning and Manipulation

This section documents the data cleaning and manipulation. It includes the following tasks:

1. Setting up and loading data into R environment
2. Removing duplicate or irrelevant observations
3. Fixing structural errors
4. Handle missing data
5. Adding new columns
6. Filter unwanted outliers

This report uses the *R programming language* in conjunction with the *RStudio* integrated development environment (IDE) to complete these tasks. All R scripts are available on the [GitHub repository](#).

3.1 Setting Up and Loading Data into R Environment

Prior to cleaning the data, it is important to set up the R environment on the project. This requires loading the following packages into the environment:

- `tidyverse` for importing and wrangling the data
- `janitor` for cleaning data
- `ggplot2` for visualizing the data
- `ggmap` for visualizing the data in a map

Once the environment was set up, each CSV file was imported into a separate data frame, and then these data frames were combined into a single data frame, named `trip_data_v2`. Upon initial inspection, the data frame contains 5,667,717 rows and 13 columns. The table below provides a preview of the `trip_data_v2` data frame.

Table 1: A preview of Cyclistic's historical trip data

ride_id	98D355D9A9852BE9
rideable_type	classic_bike
started_at	2022-01-01 00:00:05
ended_at	2022-01-01 00:01:48
start_station_name	Michigan Ave & 8th St
start_station_id	623
end_station_name	Michigan Ave & 8th St
end_station_id	623
start_lat	41.87277
start_lng	-87.62398
end_lat	41.87277
end_lng	-87.62398
member_casual	casual

3.2 Removing Duplicate or Irrelevant Observations

The first task in data cleaning and manipulation is to remove unwanted observations from the data set, including duplicate observations or irrelevant observations. Duplicate observations will happen most often during the data collection. This task involves two main parts, including removing irrelevant columns and remove duplicate values.

Initially, irrelevant columns were removed from the data set as they were not required for this analysis. This was done by using the `select()` function to remove the following columns: `start_station_id` and `end_station_id`. After that, it is important to remove duplicate records. To detect duplicate records in the `ride_id` column, the `summarize()` function was used to count the number of rows and unique rows in the data set. This function returned a data frame with two columns: `total_ride` (the number of rows) and `unique_ride` (the number of unique rows).

Table 2: The total number of rides and unique rides in Cyclistic's historical trip data

total_ride	unique_ride
5667717	5667717

The result shows that there are 5,667,717 rides in total and 5,667,717 unique rides in the data set. This implies that the data set does not contain any duplicate values. Therefore, the task of removing duplicate records can be skipped in the `ride_id` column. There are currently 5,667,717 rows and 11 columns remaining in the data set.

3.3 Fixing Structural Errors

The next task is to fix structural errors occurred in the data. This involves ensuring that all data is in a consistent formats or naming conventions, with the same data types and values. This includes bike and rider types and station names.

3.3.1 Bike and Rider Types

This step involved checking for irrelevant records in the `rideable_type` and `member_casual` columns. The `count()` function was used to count the number of records in each category.

Table 3: A list of bike type in in Cyclistic's historical trip data

rideable_type	total_ride
classic_bike	2601214
docked_bike	177474
electric_bike	2889029

Table 4: A list of user type in in Cyclistic's historical trip data

member_casual	total_ride
casual	2322032
member	3345685

Tables 3 and 4 show that there are three types of bikes: `classic_bike`, `docked_bike`, and `electric_bike`. There are also two types of riders: `casual` and `member`. These results suggest that there are no consistency errors in the `rideable_type` or `member_casual` columns.

3.3.2 Station Names

This step involved making the station names consistent in the `start_station_name` and `end_station_name` columns. The `count()` function was used again to ensure that there were no inconsistencies. The pre-cleaned versions of the `start_station_name` and `end_station_name` columns can be found in the [preclean-start-station.csv](#) and [preclean-end-station.csv](#).

However, an inspection of these columns revealed some inconsistent formats. First, the data set included test stations, which should not be included in the analysis and must be removed. There were eight test stations, including `Pawel Bialowas - Test- PBSC charging station`, `Hastings WH 2`, `DIVVY CASSETTE REPAIR MOBILE STATION`, `Base - 2132 W Hubbard Warehouse`, `Base - 2132 W Hubbard`, `NewHastings`, `WestChi`, and `WEST CHI-WATSON`. The test stations were removed using the following R code:

```
test_station_list <- c("Pawel Bialowas - Test- PBSC charging station",
                      "Hastings WH 2",
                      "DIVVY CASSETTE REPAIR MOBILE STATION",
                      "Base - 2132 W Hubbard Warehouse",
                      "Base - 2132 W Hubbard",
                      "NewHastings",
                      "WestChi",
                      "WEST CHI-WATSON")

trip_data_v2 <- trip_data_v2[
  !( trip_data_v2$start_station_name %in% test_station_list |
    trip_data_v2$end_station_name %in% test_station_list ), ]
```

Additionally, it was found that the station names contained some extraneous characters and words. These extraneous characters and words could cause problems when grouping the data by station name, so they must be removed. The

extraneous characters and words that were found include `*`, `- Charging`, `(Temp)`, and `amp;`. The following R code was used to remove these extraneous characters and words:

```
trip_data_v2 <- trip_data_v2 %>%
  mutate(start_station_name = str_replace_all(start_station_name, fixed("*"), "")) %>%
  mutate(start_station_name = str_replace_all(start_station_name, fixed("- Charging"), "")) %>%
  mutate(start_station_name = str_replace_all(start_station_name, fixed(" (Temp)"), "")) %>%
  mutate(start_station_name = str_replace_all(start_station_name, fixed("amp;"), "")) %>%
  mutate(end_station_name = str_replace_all(end_station_name, fixed("*"), "")) %>%
  mutate(end_station_name = str_replace_all(end_station_name, fixed("- Charging"), "")) %>%
  mutate(end_station_name = str_replace_all(end_station_name, fixed(" (Temp)"), "")) %>%
  mutate(end_station_name = str_replace_all(end_station_name, fixed("amp;"), ""))
```

Finally, the task of fixing structural errors has been completed. The data set now contains 5,665,349 rows and 11 columns.

3.4 Handle Missing Data

The third task in data cleaning and manipulation is to handle missing data. There are a few approaches to dealing with missing data, which include:

1. **Imputation:** This involves filling in the missing values with estimates based on other observations.
2. **Deletion:** This involves dropping the observations that have missing values.

Neither of these approaches is ideal, but both can be considered. The following table shows a list of the total number of missing values in each column.

Table 5: A list of the total number of missing values in each column
(before removing missing values)

	total_na
ride_id	0
rideable_type	0
started_at	0
ended_at	0
start_station_name	833025
end_station_name	891896
start_lat	0
start_lng	0
end_lat	5858

	total_na
end_lng	5858
member_casual	0

The table above shows that there are four columns with missing values, including:

- `start_station_name` (833,025 rows)
- `end_station_name` (891,896 rows)
- `end_lat` (5,858 rows)
- `end_lng` (5,858 rows)

In the next step, the missing values in the `start_station_name` and `end_station_name` columns will be replaced with the other observations based on geographic coordinates. The remaining missing values will be removed. The following R code was used to impute the missing station names with geographic coordinates:

```
trip_data_v2 <- trip_data_v2 %>%
  group_by(start_lat, start_lng) %>%
  fill(start_station_name, .direction = "downup") %>%
  ungroup()

trip_data_v2 <- trip_data_v2 %>%
  group_by(end_lat, end_lng) %>%
  fill(end_station_name, .direction = "downup") %>%
  ungroup()
```

The number of rows with missing values in the `start_station_name` and `end_station_name` columns is reduced to 476,854 and 559,273, respectively. The remaining missing values (including `end_lat` and `end_lng`) are then removed by running the following R code:

```
trip_data_v2 <- trip_data_v2[complete.cases(trip_data_v2), ]
```

Missing data was completely removed from the data frame. After inspecting the station names again, all missing and inconsistent values were eliminated. The post-cleaned versions of the `start_station_name` and `end_station_name` columns can also be found in the [postclean-start-station.csv](#) and [postclean-end-station.csv](#) files. As shown in the table below, there are currently no missing values in the data frame.

Table 6: A list of the total number of missing values in each column
(after removing missing values)

	total_na
ride_id	0
rideable_type	0
started_at	0
ended_at	0
start_station_name	0
end_station_name	0
start_lat	0
start_lng	0
end_lat	0
end_lng	0
member_casual	0

The task has been completed. Now, there are 4,803,698 rows and 11 columns remaining in the data set.

3.5 Adding New Columns

The fourth step in the data cleaning and manipulation process is to add three new columns to the data set:

- `month`: The month in which the trip was taken
- `day_of_week`: The day of the week on which the trip was taken
- `ride_length`: The duration of the trip in minutes

These columns were added using the `mutate()` function. It allows users to add new columns to a data set or to modify existing columns. There are currently 14 columns in the data set, as shown in Table 7. The addition of these new columns will make it possible to perform more detailed analysis of the data in the next section.

Table 7: A preview of Cyclistic's historical trip data

ride_id	98D355D9A9852BE9
rideable_type	classic_bike
started_at	2022-01-01 00:00:05
ended_at	2022-01-01 00:01:48
start_station_name	Michigan Ave & 8th St
end_station_name	Michigan Ave & 8th St
start_lat	41.87277
start_lng	-87.62398

end_lat	41.87277
end_lng	-87.62398
member_casual	casual
month	January
day_of_week	Saturday
ride_length	1.716667

3.6 Filter Unwanted Outliers

The final task in the data cleaning and manipulation process is to filter unwanted outliers. Outliers can be caused by errors in data entry, measurement, or other factors. They can distort the results of an analysis, so it is important to remove them before proceeding. In this case, all trip data that were less than 60 seconds were filtered out. These trips are likely to be false starts or users attempting to re-dock a bike to ensure its security. Trips that exceeded a period of 24 consecutive hours were also removed. This is because these trips are likely to be the result of the rider not returning the bike. This was accomplished by running the following R code:

```
trip_data_v2 <- trip_data_v2[!( trip_data_v2$ride_length < 1 |
                                trip_data_v2$ride_length > 1440), ]
```

The presence of outliers in a data set can skew the results of statistical analyses. To ensure that the data set was free of outliers, the interquartile range (IQR) method was used to filter them out. The IQR is a measure of the variability of a data set, and it is calculated by subtracting the first quartile (Q1) from the third quartile (Q3). Outliers are defined as data points that are more than 1.5 times the IQR away from Q1 or Q3. The following R code was used to remove outliers from the data set:

```
median_value <- median(trip_data_v2$ride_length)
iqr_value <- IQR(trip_data_v2$ride_length)
lower_limit <- median_value - 1.5 * iqr_value
upper_limit <- median_value + 1.5 * iqr_value

trip_data_v2 <- trip_data_v2[!( trip_data_v2$ride_length < lower_limit |
                                trip_data_v2$ride_length > upper_limit), ]
```

The trip data that were less than -8.54 minutes (technically 0 minutes) or greater than 29.91 minutes were deleted from the data set. This left 4,151,346 rows remaining after removing unwanted outliers. This completes the data cleaning and manipulation process, and the data is now ready for analysis.

4 Analysis and Key Findings

4.1 The Number of Rides

4.2 The Duration of Rides

4.3 The Location of Rides

4.4 Key Findings

Based on data analysis, :

-
-
-
-
-

5 Conclusion and Recommendations

5.1 Conclusion

Based on the data collected, it can be concluded that

5.2 Recommendations

The report makes the following recommendations:

- 1.
- 2.
- 3.